

## Exercises 05

### Stored Functions in SQL and PLpgSQL

1. Write a simple PLpgSQL function that returns the square of its argument value. It is used as follows:

```
mydb=> select sqr(4);
sqr
-----
    16
(1 row)

mydb=> select sqr(1000);
sqr
-----
1000000
(1 row)
```

Could we use this function to square real numbers? If not, how could we write a function to achieve this?

[hide answer]

```
create or replace function sqr(n integer) returns integer
as $$
declare
    result integer;
begin
    result := n * n;
    return result;
end;
$$ language plpgsql;
```

OR

```
create or replace function sqr(n integer) returns integer
as $$
begin
    return n * n;
end;
$$ language plpgsql;
```

This function won't square real numbers, even something like:

```
mydb=> select sqr(3.0);
ERROR:  Function sqr(numeric) does not exist
```

The types don't match, and PostgreSQL won't automatically convert. However, by declaring the function to use the generic number type `numeric`, it will handle both integers and reals correctly:

```
create or replace function sqr(n numeric) returns numeric
as $$
begin
    return n * n;
end;
$$ language plpgsql;
```

2. Write two PLpgSQL functions to compute factorial, one iterative and one recursive. Both functions take an integer argument and return an integer result, if the argument is non-negative. Otherwise, they return NULL. They would both be used as follows:

```
mydb=> select fac(5);
fac
-----
 120
```

```
(1 row)

mydb=> select fac(10);
      fac
-----
 3628800
(1 row)

mydb=> select fac(0);
      fac
-----
        1
(1 row)

mydb=> select fac(-1);
      fac
-----
(0 rows)
```

[hide answer]

```
-- recursive version
create or replace function fac(n integer) returns integer
as $$
begin
    if (n < 0) then
        return null;
    elsif (n = 0) then
        return 1;
    elsif (n = 1) then
        return 1;
    else
        return n * fac(n - 1);
    end if;
end;
$$ language plpgsql;

-- iterative version
create or replace function fac(n integer) returns integer
as $$
declare
    i integer;
    res integer;
begin
    if (n < 0) then
        return null;
    end if;
    res := 1;
    for i in 1 .. n loop
        res := res * i;
    end loop;
    return res;
end;
$$ language plpgsql;
```

3. Write a PLpgSQL function that "spreads" the letters in some text. It is used as follows:

```
mydb=> select spread('My Text');
      spread
-----
 M y   T e x t
(1 row)
```

[hide answer]

```
create or replace function spread(text) returns text
as $$
```

```

declare
    result text := '';
    i      integer;
    len    integer;
begin
    i := 1;
    len := length($1);
    while (i <= len) loop
        result := result || substr($1, i, 1) || ' ';
        i := i+1;
    end loop;
    return result;
end;
$$ language plpgsql;

```

OR

```

create or replace function spread(text) returns text
as $$
declare
    result text := '';
    i      integer;
begin
    i := 1;
    for i in 1..length($1) loop
        result := result || substr($1, i, 1) || ' ';
    end loop;
    return result;
end;
$$ language plpgsql;

```

Note that if you omit the initial assignment of empty string to `result`, then the value of `result` stays as `NULL` throughout the entire function execution, and `NULL` is returned (i.e., string concatenation is a null-preserving operation).

4. Write a PLpgSQL function to return a table of the first  $n$  positive integers. You may assume the existence of a tuple type:

```
create type IntValue as ( val integer );
```

The function is used as follows:

```

mydb=> select * from seq(5);
 val
-----
 1
 2
 3
 4
 5
(5 rows)

```

and has the following signature:

```
create or replace function seq(int) returns setof IntValue
```

[hide answer]

```

create or replace
    function seq(int) returns setof IntValue
as $$
declare
    i integer;
    r IntValue%rowtype;
begin
    for i in 1 .. $1
    loop

```

```

        r.val = i;
        return next r;
    end loop;
    return;
end;
$$ language plpgsql;

```

5. Generalise the previous function so that it returns a table of integers, starting from *lo* up to at most *hi*, with an increment of *inc*. The function should also be able to count down from *lo* to *hi* if the value of *inc* is negative. An *inc* value of 0 should produce an empty table. Use the following function header:

```
create or replace function seq(int,int,int) returns setof IntValue
```

and the function would be used as follows:

```

mydb=> select * from seq(2,7,2);
 val
-----
    2
    4
    6
(3 rows)

```

Some other examples, in a more compact representation:

```

seq(1,5,1)  gives  1  2  3  4  5
seq(5,1,-1) gives  5  4  3  2  1
seq(9,2,-3) gives  9  6  3
seq(2,9,-1) gives  empty
seq(1,5,0)  gives  empty

```

[hide answer]

```

create or replace function
    seq(lo int, hi int, inc int) returns setof IntValue
as $$
declare
    i integer;
    r IntVal%rowtype;
begin
    i := lo;
    if (inc > 0) then
        while (i <= hi)
        loop
            r.val = i;
            return next r;
            i := i + inc;
        end loop;
    elsif (inc < 0) then
        while (i >= hi)
        loop
            r.val = i;
            return next r;
            i := i + inc;
        end loop;
    end if;
    return;
end;
$$ language plpgsql;

```

6. Re-implement the `seq(int)` function from above as an **SQL function**, and making use of the generic `seq(int,int,int)` function defined above.

[hide answer]

```

create or replace function
    seq(n int) returns setof IntValue

```

```

as $$
declare
    r IntVal%rowtype;
begin
    for r in select * from seq(1,n,1)
    loop
        return next r;
    end loop;
    return;
end;
$$ language plpgsql;

```

or, it could be done more simply as an SQL function:

```

create or replace function
    seq(int) returns setof IntValue
as $$
    select * from seq(1,$1,1);
$$
language sql;

```

7. Create a new version of the factorial function based on the above sequence returning functions. Implement it as an **SQL function** (not a PLpgSQL function). The obvious solution to this problem requires a product aggregate, analogous to the sum aggregate. Since PostgreSQL does not have a built-in product aggregate, you will need to build your own. In order to achieve the same behaviour as the above PLpgSQL function for non-positive numbers, you will need to use the same function signature as before, i.e.:

```

create or replace function fac(int) returns int

```

Note that PostgreSQL does *not* have a product aggregate, so you will need to build your own.

[hide answer]

```

create function fac(int) returns int
as $$
    select product(val) from seq($1);
$$
language sql;

-- where the product aggregate is defined as follows:
-- basetype = type of the input values
-- stype     = type of intermediate states in computing aggregate
-- sfunc     = function mapping (currentState, nextValue) -> newState
-- initcond  = value for starting state
-- The "intermediate state" for computing a product is very simple.
-- It's simply the product accumulated so far. Similarly, the mapping
-- function simply multiplies the current product by the next value.
--
-- For more details on defining aggregates, see ...
-- PostgreSQL Reference Manual, SQL section

create aggregate product(int) (
    sfunc     = multiplyNext,
    stype     = int,
    initcond  = 1
);

create function multiplyNext(int,int) returns int
as $$
begin return $1 * $2; end;
$$ language plpgsql;

-- note that the above definitions need to be given in the reverse
-- order to what they're given here in order for this to work

```

Use the Beers/Bars/Drinkers database from lectures in answering the following questions. A summary schema for this database:

```

Beers(name:string, manufacturer:string)
Bars(name:string, address:string, license#:integer)
Drinkers(name:string, address:string, phone:string)
Likes(drinker:string, beer:string)
Sells(bar:string, beer:string, price:real)
Frequents(drinker:string, bar:string)

```

The examples below assume that the user is connected to a database called `beer` containing an instance of the above schema.

8. Write a PLpgSQL function called `hotelsIn()` that takes a single argument giving the name of a suburb, and returns the names of all hotels in that suburb, one per line. It is used as follows:

```

beer=> select hotelsIn('The Rocks');
          hotelsin
-----
Australia Hotel
Lord Nelson

(1 row)

```

Note that the output from functions returning text looks much better if you turn off output alignment (via `psql's \a` command) and get rid of column headings (via `psql's \t` command).

Compare the aligned output above to the unaligned output below:

```

beer=> \a
Output format is unaligned.
beer=> \t
Showing only tuples.
beer=> select hotelsIn('The Rocks');
Australia Hotel
Lord Nelson

```

From now on, sample outputs for functions returning text will assume that we have used `\a` and `\t`.

[hide answer]

Note that this returns a string and not a list of records.

```

create or replace function
    hotelsIn(_addr text) returns text
as $$
declare
    r    record;
    out text := '';
begin
    for r in select * from bars where addr = _addr
    loop
        out := out || r.name || e'\n';
    end loop;
    return out;
end;
$$ language plpgsql;

```

9. Write a new PLpgSQL function called `hotelsIn()` that takes a single argument giving the name of a suburb and returns the names of all hotels in that suburb. The hotel names should all appear on a single line, as in the following examples:

```

beer=> select hotelsIn('The Rocks');
Hotels in The Rocks:  Australia Hotel  Lord Nelson

beer=> select hotelsIn('Randwick');
Hotels in Randwick:  Royal Hotel

beer=> select hotelsIn('Rendwick');
Hotels in Rendwick:

```

[hide answer]

```

create or replace function
    hotelsIn (_addr text) returns text
as $$
declare
    pubnames text;
    p record;
begin
    pubnames:= 'Hotels in ' || address || ':';
    for p in select * from Bars where addr = _addr
    loop
        pubnames := pubnames||' '||p.name;
    end loop;
    pubnames := pubnames||e'\n';
    return pubnames;
end;
$$ language plpgsql;

```

10. Modify the PLpgSQL function in the previous question so that it prints a more sensible message if there are no hotels in the named suburb. It should also display the hotel names on separate lines and numbers them. It is used as follows:

```

beer=> select hotelsIn('Rendwick');
There are no hotels in Rendwick

beer=> select hotelsIn('The Rocks');
Hotels in The Rocks:
  1. Australia Hotel
  2. Lord Nelson

```

Use `to_char(i,99)` to format the numbers.

[hide answer]

This uses "select ... into" to extract a count.

```

create or replace function
    hotelsIn(_addr text) returns text
as $$
declare
    pubnames text;
    p record;
    i integer := 0;
    howmany integer := 0;
begin
    select count(*) into howmany from Bars where addr = _addr;
    if (howmany = 0) then
        return 'There are no hotels in ' || address || e'\n';
    end if;
    pubnames := 'Hotels in ' || address || e':\n';
    for p in select * from Bars where addr = _addr
    loop
        i := i+1;
        pubnames := pubnames|| to_char(i,99) ||'. '||p.name||e'\n';
    end loop;
    return pubnames;
end;
$$ language plpgsql;

```

11. Write a PLpgSQL procedure `happyHourPrice` that accepts the name of a hotel, the name of a beer and the number of dollars to deduct from the price, and returns a new price. The procedure should check for the following errors:

- non-existent hotel (invalid hotel name)
- non-existent beer (invalid beer name)
- beer not available at the specified hotel
- invalid price reduction (e.g. making reduced price negative)

Use `to_char(price, '$9.99')` to format the prices.

```
beer=> select happyHourPrice('Oz Hotel','New',0.50);
There is no hotel called 'Oz Hotel'.

beer=> select happyHourPrice('Australia Hotel','Newer',0.50);
There is no beer called 'Newer'.

beer=> select happyHourPrice('Australia Hotel','New',0.50);
The Australia Hotel does not serve New

beer=> select happyHourPrice('Australia Hotel','Burraborang Bock',4.50);
Price reduction is too large; Burraborang Bock only costs $ 3.50

beer=> select happyHourPrice('Australia Hotel','Burraborang Bock',1.50);
Happy hour price for Burraborang Bock at Australia Hotel is $ 2.00
```

[hide answer]

This could be done using either all `select count(*)` followed by a check for zero, or by using the `FOUND` variable (which is set after each query). This solution combines both approaches to show the range of possibilities.

In general, you could use `count(*)` whenever you knew that you were not interested in collecting any other information from the table; you'd try to collect the information and use `FOUND` in all other circumstances.

```
-- using positional notation for parameters
create or replace function
    happyHourPrice (_hotel text, _beer text, _discount real) returns text
as $$
declare
    counter integer;
    std_price real;
    new_price real;
begin
    select into counter count(*) from Bars where name = _hotel;
    if (counter = 0) then
        return 'There is no hotel called ' || _hotel || e'.'.n';
    end if;
    select into counter count(*) from Beers where name = _beer;
    if (counter = 0) then
        return 'There is no beer called ' || _beer || e'.'.n';
    end if;
    select price into std_price
    from Sells s
    where s.beer = _beer and s.bar = _hotel;
    if (not found) then
        return 'The ' || _hotel || ' does not serve ' || _beer;
    end if;
    new_price := std_price - _discount;
    if (new_price < 0) then
        return 'Price reduction is too large; '
            || _beer || ' only costs '
            || to_char(std_price, '$9.99');
    else
        return 'Happy hour price for '
            || _beer || ' at ' || _hotel || ' is '
            || to_char(new_price, '$9.99');
    end if;
end;
$$ language plpgsql;
```

12. The `hotelsIn` function above returns a formatted string giving details of the bars in a suburb. If we wanted to return a table of records for the bars in a suburb, we could use a view as follows:

```
beer=> create or replace view HotelsInTheRocks as
-> select * from Bars where addr = 'The Rocks';
CREATE VIEW
```



```
beer=> select * from HotelsInTheRocks;
      name      | addr      | license
-----+-----+-----
Australia Hotel | The Rocks | 123456
Lord Nelson     | The Rocks | 123888
(2 rows)
```

Unfortunately, we need to specify a suburb in the view definition. It would be more useful if we could define a "parameterised view" which we could use to generate a table for any suburb, e.g.

```
beer=> select * from HotelsIn('The Rocks');
      name      | addr      | license
-----+-----+-----
Australia Hotel | The Rocks | 123456
Lord Nelson     | The Rocks | 123888
(2 rows)
beer=> select * from hotelsIn('Coogee');
      name      | addr      | license
-----+-----+-----
Coogee Bay Hotel | Coogee    | 966500
(1 row)
```

Such a parameterised view can be implemented via an SQL function, defined as:

```
create or replace function hotelsIn(text) returns setof Bars
as $$ ... $$ language sql;
```

Complete the definition of the SQL function.

[hide answer]

```
create or replace function
    hotelsIn(text) returns setof Bars
as $$
    select * from Bars where addr = $1;
$$ language sql;
```

13. The function for the previous question can also be implemented in PLpgSQL. Give the PLpgSQL definition. It would be used in the same way as the above.

[hide answer]

```
create or replace function
    hotelsIn(_addr text) returns setof Bars
as $$
declare
    r record; -- could also be declared r Bars%rowtype;
begin
    for r in select * from Bars where addr = _addr
    loop
        return next r;
    end loop;
    return;
end;
$$ language plpgsql;
```

Use the Bank Database in answering the following questions. A summary schema for this database:

```
Branches(location:text, address:text, assets:real)
Accounts(holder:text, branch:text, balance:real)
Customers(name:text, address:text)
Employees(id:integer, name:text, salary:real)
```

The examples below assume that the user is connected to a database called bank containing an instance of the above schema.

14. For each of the following, write both an SQL and a PLpgSQL function to return the result:

## a. salary of a specified employee

[hide answer]

```
-- Salary of a specified employee
--   Allows employee to be determined by name or id
--   using overloading on the function name
--   Assume name or id identifies only one employee

create or replace function empSal(text) returns real
as $$
    select salary from employees where name = $1
$$ language sql;

create or replace function empSal(integer) returns real
as $$
    select salary from employees where id = $1
$$ language sql;

create or replace function
    empSal1(_name text) returns real
as $$
declare
    _sal real;
begin
    select salary into _sal
    from employees where name = _name;
    return _sal;
end;
$$ language plpgsql;

create or replace function
    empSal1(_id integer) returns real
as $$
declare
    _sal real;
begin
    select salary into _sal
    from employees where id = _id;
    return _sal;
end;
$$ language plpgsql;
```

## b. all details of a particular branch

[hide answer]

```
-- All details of a particular branch
--   Example of PLpgSQL function returning a record

create or replace function branchDetails(text) returns Branches
as $$
    select * from Branches where location = $1;
$$ language sql;

create or replace function branchDetails1(_bname text) returns Branches
as $$
declare
    _tup Branches;
begin
    select * into _tup
    from Branches where location = _bname;
    return _tup;
end;
$$ language plpgsql;
```

## c. names of all employees earning more than \$sal

[hide answer]

```
-- Names of all employees earning more than $sal
-- Example of PLpgSQL function returning a set of atomic values

create or replace function empsWithSal(real) returns setof text
as $$
    select name from employees where salary > $1;
$$ language sql;

create type EmpName as ( name text );

create or replace function empsWithSal1(_minSal real) returns setof EmpName
as $$
declare
    _en EmpName;
begin
    for _en in select name
                from employees where salary > _minSal
    loop
        return next _en;
    end loop;
    return;
end;
$$ language plpgsql;
```

d. all details of highly-paid employees

[hide answer]

```
-- All details of highly-paid employees
-- Example of PLpgSQL function returning a set of atomic values

create or replace function richEmps(real) returns setof Employees
as $$
    select * from employees where salary > $1;
$$ language sql;

create or replace function emps1(_minSal real) returns setof Employees
as $$
declare
    _e Employee;
begin
    for _e in select *
                from employees where salary > _minSal
    loop
        return next _e;
    end loop;
    return;
end;
$$ language plpgsql;
```

15. Write a PLpgSQL function to produce a report giving details of branches:

- name and address of branch
- list of customers who hold accounts at that branch
- total amount in accounts held at that branch

Use the following format for each branch:

```
Branch: Clovelly, Clovelly Rd.
Customers: Chuck Ian James
Total deposits: $ 8860.00
```

[hide answer]

```
create or replace function branchList() returns text
as $$
declare
```

```

a    record;
b    record;
tot integer;
qry text;
out text := e'\n';

begin
  for b in select * from Branches
  loop
    out := out || 'Branch: ' || b.location || ', ';
    out := out || b.address || e'\n' || 'Customers: ';
    tot := 0;
    for a in select * from Accounts where branch=b.location
    loop
      out := out || ' ' || a.holder;
      tot := tot + a.balance;
    end loop;
    select sum(balance) into tot
    from Accounts where branch=b.location;
    out := out || E'\nTotal deposits: ';
    out := out || to_char(tot, '$999999.99');
    out := out || E'\n---\n';
  end loop;
  return out;
end;
$$ language plpgsql;

```

It's also possible to implement this more efficiently using just one SQL query (rather than nested-loop queries). The more efficient solution involves ordering the Accounts tuples by branch, and keeping track of when the current branch changes to a new one.

Use the UNSWSIS Database in answering the following questions. Note that the UNSWSIS schema is very similar to *but not the same* as the MyMyUNSW schema (in fact, UNSWSIS is an older version of MyMyUNSW). The schema is too large to give a complete summary here, but we provide some details for some tables:

```

Person(id:integer, ..., name:text, ...)
Student(id:integer, sid:integer, stype:('local','intl'))
Staff(id:integer, sid:integer, office:integer, ...)
Term(id:integer, year:integer, session:('S1','S2','X1','X2'), ...)
Subject(id:integer, code:text, ..., name:text, ... uoc:integer, ...)
Course(id:integer, subject:integer, term:integer, lic:integer, ...)

```

The examples below assume that the user is connected to a database called `unswsis` containing an instance of the above schema.

16. Write a PLpgSQL function to produce the complete name of an OrgUnit:

```
function unitName(_oid integer) returns text
```

The function returns the complete name using the rules:

- the university is denoted by UNSW
- a faculty is denoted using its base name (not all faculty names start with Faculty)
- a school is denoted School of XYZ
- a department is denoted Department of XYZ
- a centre is denoted Centre for XYZ
- an institute is denoted Institute of XYZ
- other kinds of OrgUnits are treated as having no name (i.e. return null)

[hide answer]

```

create or replace function unitName(_oid integer) returns text
as $$
declare
  _outype text;
  _ouname text;
begin
  -- check whether the orgunit ID is valid
  select * from OrgUnit where id = _oid;

```

```

if (not found) then
    raise exception 'No such unit: %',_oid;
end if;

select t.name,u.longname into _outype,_ouname
from   OrgUnitType t, OrgUnit u
where  u.id = _oid and u.utype = t.id;

-- debugging output
-- raise notice 'Type:%, Name:%',_outype,_ouname;

if (_outype = 'UNSW') then
    return 'UNSW';
elsif (_outype = 'Faculty') then
    return _ouname;
elsif (_outype = 'School') then
    return 'School of '||_ouname;
elsif (_outype = 'Department') then
    return 'Department of '||_ouname;
elsif (_outype = 'Centre') then
    return 'Centre for '||_ouname;
elsif (_outype = 'Institute') then
    return 'Institute of '||_ouname;
else
    return null;
end if;
end;
$$ language plpgsql;

```

An alternative, using an SQL CASE expression:

```

create or replace function unitName(_oid integer) returns text
as $$
declare
    _ouname text;
begin
    -- check whether the orgunit ID is valid
    select * from OrgUnit where id = _oid;
    if (not found) then
        raise exception 'No such unit: %',_oid;
    end if;

    select case
        when t.name = 'UNSW' then 'UNSW'
        when t.name = 'Faculty' then t.longname
        when t.name = 'School' then 'School of '||t.longname
        when t.name = 'Department' then 'Department of '||t.longname
        when t.name = 'Centre' then 'Centre for '||t.longname
        when t.name = 'Institute' then 'Institute of '||t.longname
        else null
    end into _ouname
    from   OrgUnitType t, OrgUnit u
    where  u.id = _oid and u.utype = t.id;
    return _ouname;
end;
$$ language plpgsql;

```

If you didn't care about error-checking on the OrgUnit ID, then this could be done as an SQL function.

17. Write a PLpgSQL function to produce a transcript as a set of TranscriptEntry tuples, where the tuple type is defined as:

```

create type TranscriptEntry as (
    course text,      -- e.g. 'COMP1911 06s1 Computing 1'
    mark   integer,   -- e.g. 50
    grade  char(2),    -- e.g. 'PS'
);

```

```

    uoc      integer    -- e.g. 6
);

```

As well as producing the standard `TranscriptEntry` tuples, this function should also produce a "special" `TranscriptEntry` tuple at the end of each run of courses from a given session, where:

- the `course` field holds the string 'WAM for YYSN'
- the `mark` field holds the WAM for that session
- all other attributes are NULL

The WAM (weighted average mark) is an important summary of a student's achievement at UNSW. Let us assume that WAMs are computed according to the following:

- each course for which a mark is available is counted, all other courses are ignored (e.g. courses graded SY/FL have no mark and are thus not included)
- General Education is counted as a normal course for computing the WAM
- substitutions or advanced standing subjects are counted as normal courses for computing the WAM, provided that they are based on UNSW subjects
- courses with marks available are denoted `c1`, `c2`, `c3`, ...
- the units of credit for course `c1` are denoted `uoc(c1)`
- $\text{weighted\_sum} = \text{mark}(c1) * \text{uoc}(c1) + \text{mark}(c2) * \text{uoc}(c2) + \dots$
- $\text{WAM} = \text{weighted\_sum} / (\text{uoc}(c1) + \text{uoc}(c2) + \dots)$

The WAM should be computed while the course result tuples are being read; it should not be computed by invoking a new query.

[hide answer]

```

-- result tuples for internal ts() function

create type TranscriptRecord as (
    code      char(8),
    title     text,
    term      integer,
    year      integer,
    sess      char(2),
    tstart    date,
    mark      integer,
    grade     char(2),
    uoc       integer
);

-- ts() returns a list of Transcript tuples for one student

create or replace function
    ts(integer) returns setof TranscriptRecord
as $$
select s.code,s.name as title,
       t.id as term,t.year,t.sess,t.startdate as tstart,
       e.mark,e.grade,s.uoc
from   Subject s, Course c, Term t,
       CourseEnrolment e, Student stu
where  stu.sid = $1 and stu.id = e.student
       and e.course = c.id and c.term = t.id
       and c.subject = s.id
$$
language sql;

-- result tuples for transcript() function

create type TranscriptEntry as (
    course    text,      -- e.g. 'COMP1911 06s1 Computing 1'
    mark      integer,   -- e.g. 50
    grade     char(2),   -- e.g. 'PS'
    uoc       integer    -- e.g. 6

-- builds contents of special TranscriptEntry for term WAM

```

```

create or replace function
    termRec(_tr TranscriptRecord, _weighted numeric, _sumUoC integer)
    returns TranscriptEntry
as $$
declare
    _te TranscriptEntry;
begin
    _te := (null,null,null,null);
    _te.course := 'WAM for ' ||termName(_tr.year,_tr.sess);
    if (_sumUoC = 0) then
        _te.course = 'No ' ||_te.course;
    else
        _te.mark = (_weighted/_sumUoC)::integer;
    end if;
    _te.course := '== ' ||_te.course;
    return _te;
end;
$$ language plpgsql;

-- main transcript() function that drives everything else

create or replace function
    transcript(_sid integer) returns setof TranscriptEntry
as $$
declare
    _stu      Student;
    _termWeighted numeric := 0;
    _termSumUoC integer := 0;
    _tr       TranscriptRecord;
    _prev     TranscriptRecord;
    _te       TranscriptEntry;
begin
    select * into _stu from Student where sid = _sid;
    if (not found) then
        raise exception 'No such student: %',_sid;
    end if;

    -- foreach transcript record, in the correct order

    for _tr in select * from ts(_sid) order by tstart,code
    loop
        if (_prev.term <> _tr.term) then
            -- transition from one term to next
            -- summarise any existing term info
            if (_prev.term is not null) then
                _te := termRec(_prev,_termWeighted,_termSumUoC);
                return next _te;
            end if;
            -- set up to compute wam for next term
            _termWeighted := 0;
            _termSumUoC := 0;
        end if;
        _te.course := _tr.code||' '||termName(_tr.year,_tr.sess)
            ||' '||_tr.title;
        _te.mark := _tr.mark;
        _te.grade := _tr.grade;
        _te.uoc := _tr.uoc;
        return next _te;
        if (_tr.mark is not null) then
            _termWeighted := _termWeighted + (_tr.mark * _tr.uoc);
            _termSumUoC := _termSumUoC + _tr.uoc;
        end if;
        _prev := _tr;
    end loop;
    if (_prev.term is not null) then
        _te := termRec(_prev,_termWeighted,_termSumUoC);
        return next _te;
    end if;
end;

```

```

        return;
    end;
    $$ language plpgsql;

```

18. Write a PLpgSQL function which takes the numeric identifier of a given OrgUnit and returns the numeric identifier of the parent faculty for the specified OrgUnit:

```
function facultyOf(_oid integer) returns integer
```

Note that a faculty is treated as its own parent. Note also that some OrgUnits don't belong to any faculty; such OrgUnits should return a null result from the function.

[hide answer]

```

create or replace function facultyOf(_oid integer) returns integer
as $$
declare
    _count integer;
    _tname text;
    _parent integer;
begin
    select count(*) into _count
    from OrgUnit where id = _oid;
    if (_count = 0) then
        raise exception 'No such unit: %', _oid;
    end if;

    select t.name into _tname
    from OrgUnit u, OrgUnitType t
    where u.id = _oid and u.utype = t.id;

    if (_tname is null) then
        return null;
    elsif (_tname = 'University') then
        return null;
    elsif (_tname = 'Faculty') then
        return _oid;
    else
        select owner into _parent
        from UnitGroups where member = _oid;
        return facultyOf(_parent);
    end if;
end;
$$ language plpgsql;

```

An alternative way of checking the existence of the specified organisational unit would be:

```

select * from OrgUnit where id = _oid;
if (not found) then
    raise exception 'No such unit: %', _oid;
end if;

```