

Exercises 09

Transaction Processing, Concurrency Control

1. [based on Ramakrishnan, ex.17.1]

Give a brief definition for each of the following terms:

- transaction
- dirty read
- serializable schedule
- conflict-serializable schedule
- view-serializable schedule
- two-phase locking protocol

[hide answer]

a. *transaction*

An execution of a user program that performs some action that is treated as atomic according to the semantics of some database application. The DBMS sees the transaction as a sequence of actions that can include read and write operations on the database, as well as computations.

b. *dirty read*

When a transaction reads an object that has been modified by another not-yet-committed transaction.

c. *serializable schedule*

A schedule over a set of transactions that produces a result that is the same as some serial execution of the transactions.

d. *conflict-serializable schedule*

A schedule is conflict-serializable if it is conflict-equivalent to some serial schedule. Two schedules are conflict-equivalent if they involve the same set of actions and they order every pair of conflicting actions in the same way.

e. *view-serializable schedule*

A schedule is view-serializable if it is view-equivalent to some serial schedule. Two schedules are view-equivalent if they satisfy:

- the initial value of any object is read by the same transaction in both schedules, and
- the final value of any object is written by the same transaction in both schedules, and
- any shared object is written-then-read by the same pair of transactions in both schedules.

f. *two-phase locking protocol*

The two-phase locking protocol is a way of ensuring that only serializable schedules occur when transactions execute concurrently. Under this protocol each transaction must:

- acquire a shared lock on an object before reading it
- acquire an exclusive lock on an object before writing it
- not acquire any new locks once it has released a lock

2. Draw the precedence graph for the following schedule (where C means "commit"):

T1:	R(A)	W(Z)		C
T2:		R(B)	W(Y)	C
T3:	W(A)		W(B)	C

[hide answer]

It has an edge from T3 to T1 (because of A) and an edge from T2 to T3 because of B.

This gives: T2 \rightarrow T3 \rightarrow T1

3. [based on Ramakrishnan, ex.17.2]

Consider the following incomplete schedule S:

T1:	R(X)	R(Y)	W(X)		W(X)
T2:			R(Y)		R(Y)
T3:				W(Y)	

- Determine (by using a precedence graph) whether the schedule is serializable
- Modify S to create a complete schedule that is conflict-serializable

[hide answer]

- Determine (by using a precedence graph) whether the schedule is serializable*

The precedence graph has an edge, from T1 to T3, because of the conflict between T1:R(Y) and T3:W(Y). It also has an edge, from T2 to T3, because of the conflict between the first T2:R(Y) and T3:W(Y). It also has an edge, from T3 to T2, because of the conflict between T3:W(Y) and the second T2:R(Y).

- Modify S to create a complete schedule that is conflict-serializable*

Trick question. It is not possible. Since the precedence graph is cyclic, we know that it's not conflict-serializable.

If we are allowed to add **abort** actions (which was not mentioned in the question), we could simply abort either T2 or T3 and the schedule would become conflict-serializable.

4. [based on Ramakrishnan, ex.17.3]

For each of the following schedules, state whether it is conflict-serializable and/or view-serializable. If you cannot decide whether a schedule belongs to either class, explain briefly. The actions are listed in the order they are scheduled, and prefixed with the transaction name.

- T1:R(X) T2:R(X) T1:W(X) T2:W(X)
- T1:W(X) T2:R(Y) T1:R(Y) T2:R(X)
- T1:R(X) T2:R(Y) T3:W(X) T2:R(X) T1:R(Y)
- T1:R(X) T1:R(Y) T1:W(X) T2:R(Y) T3:W(Y) T1:W(X) T2:R(Y)
- T1:R(X) T2:W(X) T1:W(X) T3:W(X)

[hide answer]

The methods used to determine these solutions:

- for conflict-serializability, draw precedence graph and look for cycles
- for view-serializability, apply the definition from lecture notes.

You can short-circuit the view serializability check. As soon as you know that the schedule is conflict-serializable, it must also be view serializable.

Solutions:

- not conflict-serializable, not view-serializable
- conflict-serializable, view-serializable
- conflict-serializable, view-serializable

d. not conflict-serializable, not view-serializable

e. not conflict-serializable, view-serializable (view equivalent to the serial schedule T1, T2, T3)

5. Is the following schedule serializable? Show your working.

T1:	R(X)W(X)W(Z)	R(Y)W(Y)
T2:	R(Y)W(Y)R(Y)	W(Y)R(X)W(X)R(V)W(V)

[hide answer]

When we talk about serializability and don't specifically say what kind, we usually mean conflict-serializable. As above, the "working" for this question involves constructing a precedence graph, based on conflicting operations, and looking for cycles.

In this case there's a conflict between T1:R(X) and T2:W(X), giving a graph edge from T1 to T2. There's also a conflict between T2:R(Y) and T1:W(Y), giving a graph edge from T2 to T1. This means the graph has a cycle, so the schedule is not serializable.

6. Consider the following two transactions:

T1	T2
-----	-----
read(A)	read(B)
A := 10*A+4	B := 2*B+3
write(A)	write(B)
read(B)	read(A)
B := 3*B	A := 100-A
write(B)	write(A)

- Write versions of the above two transactions that use two-phase locking.
- Is there a non-serial schedule for T1 and T2 that is serializable? Why?
- Can a schedule for T1 and T2 result in deadlock? If so, give an example schedule. If not, explain why not.

[hide answer]

a. *Write versions of the above two transactions that use two-phase locking.*

The basic idea behind two-phase locking is that you take out all the locks you need, do the processing, and then release the locks. Thus two-phase implementations of T1 and T2 would be:

T1	T2
-----	-----
write_lock(A)	write_lock(B)
read(A)	read(B)
A := 10*A+4	B := 2*B+3
write(A)	write(B)
write_lock(B)	write_lock(A)
read(B)	read(A)
B := 3*B	A := 100-A
write(B)	write(A)
unlock(A)	unlock(B)
unlock(B)	unlock(A)

b. *Is there a non-serial schedule for T1 and T2 that is serializable? Why?*

No. There is no such schedule. The last operation in T1 is write(B), and the last operation in T2 is write(A). T1 starts with read(A) and T2 starts with read(B). Therefore, in any serializable schedule, we would require that either read(A) in T1 should be after write(B) in T2 or read(B) in T2 should be after write(B) in T1.

c. Can a schedule for $T1$ and $T2$ result in deadlock? If so, give an example schedule. If not, explain why not.

Yes. Consider the following (where $L(X)$ denotes taking a lock on object X):

$T1:$	$L(A)R(A)$	$W(A)L(B)$	\dots
$T2:$	$L(B)$	$R(B)W(B)L(A)$	\dots